

Assignment Project Exam Help

Disk Management

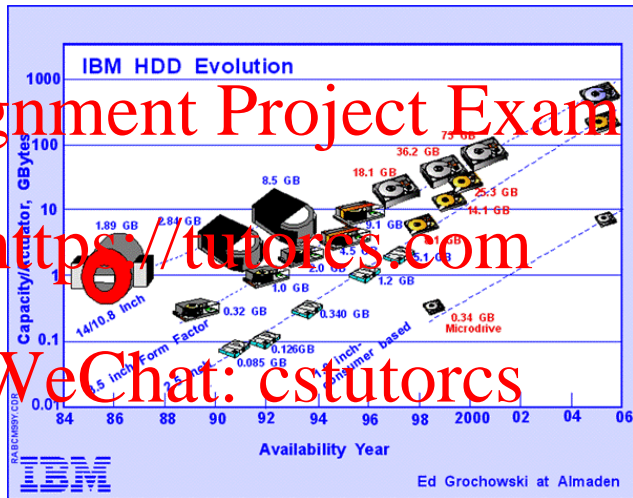
Anandha Gopalan

(with thanks to D. Rueckert, P. Pietzuch, A. Tannenbaum and
R. Kolcun)

<https://tutorcs.com>

axgopala@imperial.ac.uk

WeChat: cstutorcs



Capacity increases exponentially, but access speeds not so much

Assignment Project Exam Help

<https://tutorcs.com>



WeChat: cstutorcs

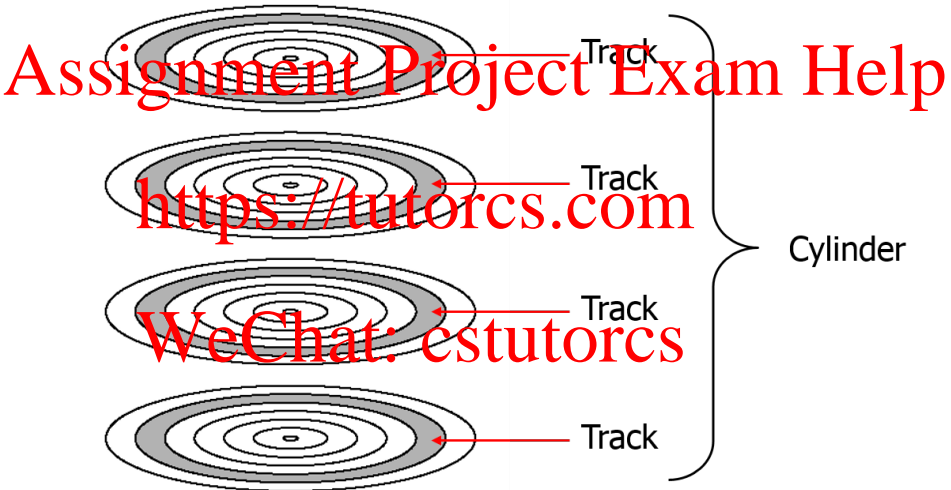


Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs





Sample Disk Specification




Assignment Project Exam Help

Parameter	IBM 360 KB floppy disk	Seagate Barracuda ST3400832AS
No. of cylinders	40	16,383
Tracks/cylinder	2	16
Sectors/track	9	63
Bytes/sector	512	512
Sectors/disk	720	781,422,768
Disk capacity	360 KB	400 GB


<http://disctech.com/Seagate-ST3400832AS-SATA-Hard-Drive>

Assignment Project Exam Help

<https://tutorcs.com>



(a) Physical geometry



(b) Virtual geometry

WeChat: cstutorcs

Surface divided into 20 or more **zones**

- Outer zones have more sectors per track → ensures that sectors have same physical length
- Zones hidden using virtual geometry

Physical hardware address: (cylinder, surface, sector)

- But actual geometry complicated → hide from OS

Modern disks use logical sector addressing (or logical block addresses LBA)

- Sectors numbered consecutively from 0 ... n
- Makes disk management much easier
- Helps work around BIOS limitations
 - Original IBM PC BIOS → 8 GB max
 - 6 bits for sector, 4 bits for head, 14 bits for cylinder

Disk capacity statements can be confusing ☹

Assignment Project Exam Help

1 KB = 2^{10} bytes = 1024 bytes vs 1 KB = 10^3 bytes = 1000 bytes

1 MB = 2^{20} bytes = 1024² bytes vs 1 MB = 10^6 bytes = 1000² bytes

1 GB = 2^{30} bytes = 1024³ bytes vs 1 GB = 10^9 bytes = 1000³ bytes

<https://tutorcs.com>
WeChat: cstutorcs

If necessary, just make it consistent on the exam ☺

Before a disk can be used, it must be formatted

- **Low level format**

- Disk sector layout



- Cylinder skew

- Interleaving

- **High level format**

- Boot block
 - Free block list
 - Root directory
 - Empty file system

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Amount of cylinder skew depends on the drive geometry

Example Problem

Consider a 10,000 rpm drive with each track having 300 sectors and track to track seek time of 800 μsec

Time taken for 1 rotation = $\frac{60\text{s}}{10000} = 6 \times 10^{-3} = 6 \text{ ms}$

300 sectors per track \Rightarrow Time taken for 1 sector = $\frac{6\text{ms}}{300} = 2 \times 10^{-5} = 20 \mu\text{s}$


Track to track seek time is 800 $\mu\text{s} \Rightarrow$ Number of sectors that pass in one seek = $\frac{800}{20} = 40$

Hence, cylinder skew = 40

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Typical disk

Sector size	512 bytes
Seek time (adjacent cylinder)	< 1 ms
Seek time (average)	8 ms
Rotation time (average latency)	4 ms
Transfer rate	up to 100 MB/s

Disk Scheduling

- Minimise seek and/or latency times
- Order pending disk requests with respect to head position

Seek time $\approx 2 - 3$ times larger than latency time \rightarrow more important to optimise

Given

b – number of bytes to be transferred

N – number of bytes per track

r – rotation speed in revolutions per second

Seek time

t_{seek}

Latency time (rotational delay)

$$t_{latency} = \frac{1}{2 \times r}$$

Transfer time

$$t_{transfer} = \frac{b}{N \times r}$$

Total access time (t_{access})

$$t_{seek} + t_{latency} + t_{transfer}$$

<https://tutorcs.com>

WeChat: cstutorcs

Disk Performance

Average seek time: 10ms

Rotation speed: 10,000 rpm

512 byte sectors

320 sectors per track

File size: 2560 sectors (1.25 MB)

Calculate the time taken to:

- 1 read file stored as compactly as possible on disk (i.e. file occupies all sectors on 8 adjacent tracks $\rightarrow 8 \text{ tracks} \times 320 \text{ sectors/track} = 2560 \text{ sectors}$)
- 2 read file with all sectors randomly distributed across disk

Example Problem

Answer: Disk Performance

$$\text{Average seek} = 10 \text{ ms}$$

$$\text{Latency time} = 3 \text{ ms} = \frac{1}{2 \times \left(\frac{10000}{60}\right)}$$

$$\text{Read 320 sectors} = 6 \text{ ms} = \frac{b}{N \times \left(\frac{10000}{60}\right)}$$

$$\text{Total (for first track)} = 19 \text{ ms}$$

$$\text{Time to read next track} = 3 \text{ ms} + 6 \text{ ms} = 9 \text{ ms}$$

$$\text{Total time} = 19 \text{ ms} + 7 \times 9 \text{ ms} = 82 \text{ ms} = 0.082 \text{ seconds}$$

Average seek and latency time same as above

$$\text{Read 1 sector} = 0.01875 \text{ ms} = \frac{512}{512 \times 320 \times \left(\frac{10000}{60}\right)}$$

$$\text{Total} = 13.01875 \text{ ms}$$

$$\text{Total time} = 2560 \times 13.01875 \text{ ms} = 33.328 \text{ seconds}$$

First Come First Served (FCFS)

No ordering of requests \rightarrow random seek patterns

- OK for lightly-loaded disks
- But poor performance for heavy loads
- Fair scheduling

Assignment Project Exam Help

Queue: 98, 183, 37, 122, 14, 130, 60, 67 (head starts at 53)

<https://tutorcs.com>

WeChat: estutorcs



Shortest Seek Time First (SSTF)

Order requests according to shortest seek distance from current head position

- Discriminates against innermost/outermost tracks

- Unpredictable and unfair performance

Queue: 98, 183, 37, 122, 14, 130, 60, 67 (head starts at 53)

<https://tutorcs.com>

WeChat: cstutorcs

If, when handling request at 14, new requests arrive for 50, 70, 100 → long delay before 183 serviced


SCAN Scheduling

Choose requests which result in shortest seek time in preferred direction

- Only change direction when reaching outermost/innermost cylinder (or no further requests in preferred direction)
- Most common scheduling algorithm (also called elevator scheduling)
- Long delays for requests at extreme locations

Queue: 98, 183, 37, 122, 14, 130, 60, 67 (head starts at 53 and direction is towards 0)

WeChat: cstutorcs



C-SCAN

Services requests in one direction only

- When head reaches innermost request, jump to outermost request
- Lower variance of requests on extreme tracks
- May delay requests indefinitely (though less likely)

Queue: 98, 183, 37, 122, 14, 130, 60, 67 (head starts at 53)

<https://tutorcs.com>

WeChat: cstutorcs



N-Step SCAN

As for SCAN, but services only requests waiting when sweep began

- Requests arriving during sweep serviced during return sweep
- Doesn't delay requests indefinitely

Queue: 98, 133, 37, 122, 14, 130, 60, 67 (head starts at 53; direction \rightarrow 0); requests 80, 140 arrive when head moving outwards

<https://tutorcs.com>

WeChat: cstutorcs



I/O requests placed in request list

- One request list for each device in system
- **bio** structure: associates memory pages with requests

Block device drivers define request operation called by kernel

- Kernel passes ordered request list
- Driver must perform all operations in list
- Device drivers do not define read/write operations

Some devices drivers (e.g. RAID) order their own requests

- Bypass kernel for request list ordering

Default: variation of SCAN algorithm

- Kernel attempts to merge requests to adjacent blocks
- But, synchronous read requests may starve during large writes

Deadline scheduler: ensures reads performed by deadline

- Eliminates read request starvation

Anticipatory scheduler: delay after read request completes

- Idea: process will issue another synchronous read operation before its quantum expires
- Reduces excessive seeking behaviour
- Can lead to reduced throughput if process does not issue another read request to nearby location
 - Anticipate process behaviour from past behaviour

Problem

- CPU performance doubling every 18 months
- Disk performance has increased only 10 times since 1970

Assignment Project Exam Help

Solution

- Use parallel disk I/O → appears to OS as a single disk

<https://tutorcs.com>

RAID (Redundant Array of Inexpensive Disks)

- Array of physical drives appearing as single virtual drive
- Stores data distributed over array of physical disks to allow parallel operation (called striping)

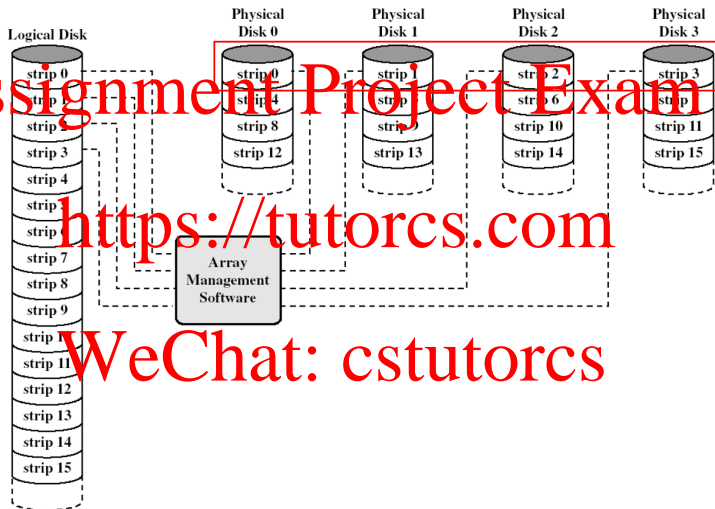
WeChat: cstutores

Use redundant disk capacity to respond to disk failure

- More disks → lower mean-time-to-failure (MTTF)

RAID

Striping



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

RAID

Level 0 (Striping)


Use multiple disks and spread out data

Disks can seek/transfer data concurrently

- May also balance load across disks

No redundancy → no fault tolerance

<https://tutorcs.com>



RAID

Level 1 (Mirroring)


Mirror data across disks

Reads can be serviced by either disk (fast)

Writes update both disks in parallel (slower)

Failure recovery easy

High storage overhead (high cost)



RAID

Level 2 (Bit-level Hamming)

Parallel access by striping at bit-level


- Use Hamming error-correcting code (ECC)
- Corrects single bit errors (and detect double-bit errors)

Very high throughput for reads/writes

- But all disks participate in I/O requests (no concurrency)
- Read-modify-write cycle

Only used if high error rates expected

- ECC disks become bottleneck
- High storage overhead



RAID


Level 3 (Byte-level XOR)

Only single parity strip used

- Parity = $\text{data1} \oplus \text{data2} \oplus \text{data3} \dots$
- Reconstruct missing data from parity and remaining data

Lower storage overhead than RAID Level 2

- But still only one I/O request can take place at a time



RAID

Level 4 (Block-level XOR)


Parity strip handled on block basis

- Each disk operates independently

Potential to service multiple reads concurrently

Parity disk tends to become bottleneck

- Data and parity strips must be updated on each write



RAID

Level 5 (Block-level Distributed XOR)


Like RAID Level 4, but distribute parity

- Most commonly used

Some potential for write concurrency

Good storage efficiency/redundancy trade-off

- Reconstruction of failed disk non-trivial (and slow)



<https://tutorcs.com>

WeChat: cstutorcs

RAID

Summary

Category	Level	Description	I/O Data Transfer (R/W)	I/O Request rate (R/W)
Striping	0	Non-redundant	+/+	+/+
Mirroring	1	Mirrored	+/0	+/0
Parallel access	2	Redundant via Hamming code	++/++	0/0
	3	Bit interleaved parity	++/++	0/0
Independent access	4	Block interleaved parity	+/-	+/-
	5	Block interleaved distributed parity	+/-	+/- or 0

better than single disk (+) / same (0) / worse (-)

Assignment Project Exam Help

Use Piazza for Q & A

Marked coursework will be returned in January

Provide feedback on PGSQL E@
<https://tutorcs.com>

Feedback also possible through Mentimeter (94 41 03)

If time permits, possible guest lecture on Security
WeChat: cstutorcs